

```


# Step 1: Upload Dataset (Only for Google Colab)
from google.colab import files
uploaded = files.upload()
# Step 2: Load the CSV
import pandas as pd
df = pd.read_csv("/content/drive/MyDrive/elonmusk.csv")
# Replace with actual filename if different
df.head()
# Step 3: Basic Info
print("Shape:", df.shape)
print("Columns:", df.columns.tolist())
df.info()
print(df.describe())
# Step 4: Missing values and duplicates
print("Missing values:\n", df.isnull().sum())
print("Duplicate rows:", df.duplicated().sum())
# Step 5: Visualizations
import seaborn as sns
import matplotlib.pyplot as plt
df['tweet_length'] = df['tweet'].astype(str).apply(len)
# Tweet length distribution
sns.histplot(df['tweet_length'], kde=True, bins=30, color='teal')
plt.title('Distribution of Tweet Lengths')
plt.xlabel('Tweet Length')
plt.ylabel('Frequency')
plt.show()
# Tweets per year
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
# Step 1: Convert 'date' column to datetime format
df['date'] = pd.to_datetime(df['date'], errors='coerce')
# Step 2: Extract year from 'date'
df['year'] = df['date'].dt.year
# Step 3: Plot number of tweets per year
sns.countplot(data=df, x='year', palette='coolwarm')
plt.title('Number of Tweets per Year')
plt.xlabel('Year')
plt.ylabel('Tweet Count')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
# Top usernames
top_users = df['username'].value_counts().head(10)
top_users.plot(kind='bar', color='orange')
plt.title('Top 10 Users by Number of Tweets')
plt.xlabel('Username')
plt.ylabel('Tweet Count')
plt.show()
# Step 6: Sentiment Analysis using TextBlob
!pip install -q textblob
from textblob import TextBlob
def get_sentiment(text):
    polarity = TextBlob(str(text)).sentiment.polarity
    if polarity > 0:
        return 'positive'
    elif polarity == 0:
        return 'neutral'
    else:
        return 'negative'
df['sentiment'] = df['tweet'].apply(get_sentiment)
print(df[['tweet', 'sentiment']].head())
# Step 7: Train-Test Split
from sklearn.model_selection import train_test_split
X = df['tweet'].astype(str)
y = df['sentiment']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
stratify=y, random_state=42)
# Step 8: TF-IDF + Logistic Regression Pipeline
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
model = Pipeline([
    ('tfidf', TfidfVectorizer(max_features=5000, stop_words='english')),
    ('clf', LogisticRegression(max_iter=1000))
])

```

```

])
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
# Step 9: Evaluation
print("🟢 Accuracy:", accuracy_score(y_test, y_pred))
print("📊 Classification Report:\n", classification_report(y_test, y_pred))
# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
# Step 10: Test on new tweets
new_tweets = [
    "Tesla's new update is amazing!",
    "I'm not happy with Twitter's new algorithm.",
    "SpaceX launch was a huge success!",
    "This is just disappointing."
]
predictions = model.predict(new_tweets)
for tweet, sentiment in zip(new_tweets, predictions):
    print(f"Tweet: {tweet}\nPredicted Sentiment: {sentiment}\n")
# Step 11: Gradio Web App
!pip install -q gradio
import gradio as gr
def predict_sentiment(tweet):
    return model.predict([tweet])[0]
iface = gr.Interface(
    fn=predict_sentiment,
    inputs=gr.Textbox(lines=3, placeholder="Enter a tweet..."),
    outputs="text",
    title="Elon Musk Tweet Sentiment Analyzer",
    description="Enter a tweet related to Elon Musk to classify sentiment as Positive, Neutral, or Negative."
)
iface.launch()

```

 Choose Files

elonmusk.csv (1).zip

- **elonmusk.csv (1).zip**(application/x-zip-compressed) - 977036 bytes, last modified: 5/19/2025 - 100% done

Saving elonmusk.csv (1).zip to elonmusk.csv (1).zip

Shape: (9286, 34)

Columns: ['id', 'conversation_id', 'created_at', 'date', 'time', 'timezone', 'user_id', 'username', 'name', 'place', 'tweet', 'mentions']

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 9286 entries, 0 to 9285

Data columns (total 34 columns):

#	Column	Non-Null Count	Dtype
0	id	9286 non-null	int64
1	conversation_id	9286 non-null	int64
2	created_at	9286 non-null	int64
3	date	9286 non-null	object
4	time	9286 non-null	object
5	timezone	9286 non-null	object
6	user_id	9286 non-null	int64
7	username	9286 non-null	object
8	name	9286 non-null	object
9	place	0 non-null	float64
10	tweet	9286 non-null	object
11	mentions	9286 non-null	object
12	urls	9286 non-null	object
13	photos	9286 non-null	object
14	replies_count	9286 non-null	int64
15	retweets_count	9286 non-null	int64
16	likes_count	9286 non-null	int64
17	hashtags	9286 non-null	object
18	cashtags	9286 non-null	object
19	link	9286 non-null	object
20	retweet	9286 non-null	bool
21	quote_url	355 non-null	object
22	video	9286 non-null	int64
23	near	0 non-null	float64
24	geo	0 non-null	float64
25	source	0 non-null	float64
26	user_rt_id	0 non-null	float64
27	user_rt	0 non-null	float64
28	retweet_id	0 non-null	float64
29	reply_to	9286 non-null	object
30	retweet_date	0 non-null	float64
31	translate	0 non-null	float64
32	trans_src	0 non-null	float64
33	trans_dest	0 non-null	float64

dtypes: bool(1), float64(11), int64(8), object(14)

memory usage: 2.3+ MB


	id	conversation_id	created_at	user_id	place	\
count	9.286000e+03	9.286000e+03	9.286000e+03	9286.0	0.0	
mean	1.055061e+18	1.052389e+18	1.540381e+12	44196397.0	NaN	
std	1.695110e+17	1.740050e+17	4.041457e+10	0.0	NaN	
min	5.610022e+17	1.659576e+09	1.422588e+12	44196397.0	NaN	
25%	9.667040e+17	9.610311e+17	1.519315e+12	44196397.0	NaN	
50%	1.089662e+18	1.088672e+18	1.548631e+12	44196397.0	NaN	
75%	1.187852e+18	1.187320e+18	1.572041e+12	44196397.0	NaN	
max	1.282940e+18	1.282933e+18	1.594712e+12	44196397.0	NaN	

	replies_count	retweets_count	likes_count	video	near	geo	\
count	9286.000000	9286.000000	9.286000e+03	9286.000000	0.0	0.0	
mean	512.958432	2282.844066	1.873480e+04	0.007538	NaN	NaN	
std	1720.320306	10546.775719	5.929579e+04	0.086500	NaN	NaN	
min	0.000000	0.000000	2.100000e+01	0.000000	NaN	NaN	
25%	43.250000	43.000000	9.430000e+02	0.000000	NaN	NaN	
50%	103.000000	129.000000	2.341500e+03	0.000000	NaN	NaN	
75%	371.000000	955.750000	1.159850e+04	0.000000	NaN	NaN	
max	49529.000000	384289.000000	1.682551e+06	1.000000	NaN	NaN	

	source	user_rt_id	user_rt	retweet_id	retweet_date	translate	\
count	0.0	0.0	0.0	0.0	0.0	0.0	
mean	NaN	NaN	NaN	NaN	NaN	NaN	
std	NaN	NaN	NaN	NaN	NaN	NaN	
min	NaN	NaN	NaN	NaN	NaN	NaN	
25%	NaN	NaN	NaN	NaN	NaN	NaN	
50%	NaN	NaN	NaN	NaN	NaN	NaN	
75%	NaN	NaN	NaN	NaN	NaN	NaN	
max	NaN	NaN	NaN	NaN	NaN	NaN	

	trans_src	trans_dest
count	0.0	0.0
mean	NaN	NaN
std	NaN	NaN
min	NaN	NaN
25%	NaN	NaN
50%	NaN	NaN
75%	NaN	NaN

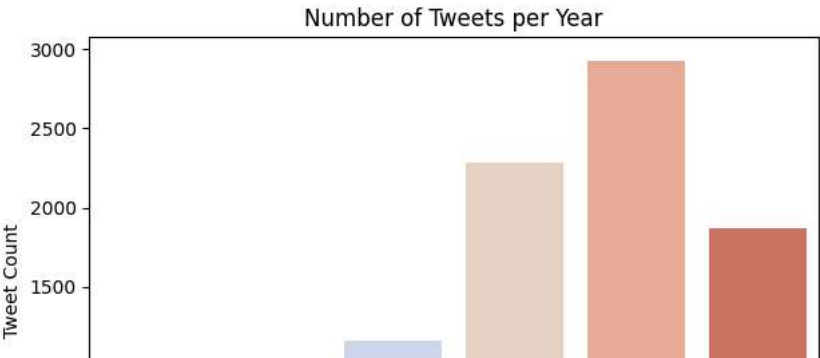
```
max      NaN      NaN
Missing values:
  id      0
  conversation_id  0
  created_at  0
  date      0
  time      0
  timezone  0
  user_id    0
  username    0
  name      0
  place      9286
  tweet      0
  mentions    0
  urls      0
  photos      0
  replies_count  0
  retweets_count  0
  likes_count  0
  hashtags    0
  cashtags    0
  link      0
  retweet     0
  quote_url  8931
  video      0
  near      9286
  geo      9286
  source     9286
  user_rt_id  9286
  user_rt    9286
  retweet_id  9286
  reply_to    0
  retweet_date  9286
  translate  9286
  trans_src  9286
  trans_dest  9286
dtype: int64
Duplicate rows: 0
```

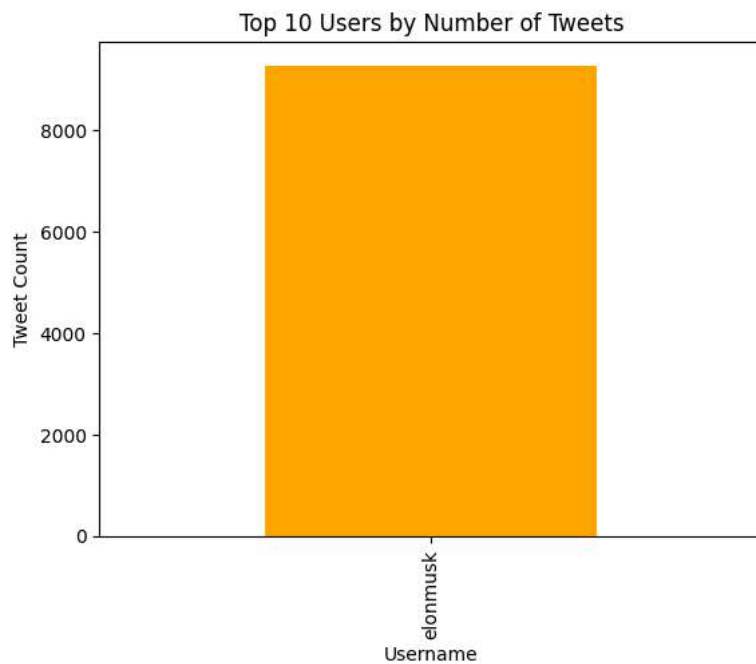



```
<ipython-input-23-f4b6e97546a9>:36: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

```
sns.countplot(data=df, x='year', palette='coolwarm')
```





```


0          tweet sentiment
1          Cute           positive
2          Wow            positive
3  Reusability is essential. A rocket that is sin... negative
4          Wild times!    positive
5  We're being extra paranoid. Maximizing probabi... positive

```

✓ Accuracy: 0.7981700753498385

📊 Classification Report:

	precision	recall	f1-score	support
negative	0.87	0.33	0.48	247
neutral	0.76	0.90	0.83	758
positive	0.83	0.85	0.84	853
accuracy			0.80	1858
macro avg	0.82	0.69	0.71	1858
weighted avg	0.81	0.80	0.78	1858



Tweet: Tesla's new update is amazing!
Predicted Sentiment: positive

Tweet: I'm not happy with Twitter's new algorithm.
Predicted Sentiment: positive

Tweet: SpaceX launch was a huge success!
Predicted Sentiment: positive

Tweet: This is just disappointing.
Predicted Sentiment: positive

It looks like you are running Gradio on a hosted a Jupyter notebook. For the Gradio app to work, sharing must be enabled. Automatically

Colab notebook detected. To show errors in colab notebook, set debug=True in launch()

* Running on public URL: <https://3d8f9c263fe9b4863d.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working dir

Elon Musk Tweet Sentiment Analyzer

Enter a tweet related to Elon Musk to classify sentiment as Positive, Neutral, or Negative.

tweet

i feel very bad

output

negative